

UC San Diego

DSC 102

Systems for Scalable Analytics

Rod Albuyeh

Midterm Review

Admin


PA0 Debrief

- Average score: 7.18/8
- Five teams claimed runtime extra credit
- New runtime record: 241.73 seconds

Midterm Overview

Similar format to sample midterms, lighter on the binary + hexadecimal questions and heavier on cloud.

Recap: Memory Hierarchy



Recap: Digital Representation of Data

Q: How many unique data items can be represented by 3 bytes?

- ❖ Given k bits, we can represent 2^k unique data items
- ❖ 3 bytes = 24 bits $\Rightarrow 2^{24}$ items, i.e., 16,777,216 items
- ❖ Common approximation: 2^{10} (i.e., 1024) $\sim 10^3$ (i.e., 1000); kibibyte (KiB) = 1024 bytes, vs kilobyte (KB) = 1000 bytes

Q: How many bits are needed to distinguish 97 unique items?

- ❖ For k unique items, invert the exponent to get $\log_2(k)$
- ❖ But #bits is an integer! So, we only need $\lceil \log_2(k) \rceil$
- ❖ So, we only need the next higher power of 2
- ❖ So... 7 bits

Recap: Decimal <-> Binary

Q: How to convert from decimal to binary representation?

1. Given decimal n

if n is power of 2 (say, 2^k), put 1 at bit position k; if k=0, stop; else pad with trailing 0s till position 0

if n is not power of 2, identify the power of 2 just below n (say, 2^k); #bits is then k; put 1 at position k

2. Reset n as $n - 2^k$; return to Steps 1-2

3. Fill remaining positions in between with 0s

Decimal	7	6	5	4	3	2	1	0	Position/Exponent of 2
	128	64	32	16	8	4	2	1	Power of 2
5_{10}						1	0	1	
47_{10}				1	0	1	1	1	
163_{10}	1	0	1	0	0	0	1	1	
16_{10}				1	0	0	0	0	

Q: Binary to decimal?

Recap: Hexadecimal representation

- ❖ *Hexadecimal* representation is a common stand-in for binary representation; more succinct and readable
 - ❖ Base 16 instead of base 2 cuts display length by ~4x
 - ❖ Digits are 0, 1, ... 9, A (10_{10}), B, ... F (15_{10})
 - ❖ Each hexadecimal digit represents 4 bits.

Decimal	Binary	Hexadecimal	
5_{10}	101_2	5_{16}	Alternative notations
47_{10}	$10\ 1111_2$	$2F_{16}$	
163_{10}	$1010\ 0011_2$	$A3_{16}$	$0xA3$ or $A3H$
16_{10}	$1\ 0000_2$	$1\ 0_{16}$	

Hexadecimal representation continued

Let's unpack:

Base 10...

0 1 2 3 4 5 6 7 8 9

Base 2...

0 1

Base-16 Hexadecimal...

0 1 2 3 4 5 6 7 8 9 A B C D E F
10 11 12 13 14 15



An aside: Hexadecimal to binary relationship

8	4	2	1	Hex. Rep.
0	0	0	0	0
0	0	0	1	1
0	0	1	0	2
0	0	1	1	3
0	1	0	0	4
0	1	0	1	5
0	1	1	0	6
0	1	1	1	7
1	0	0	0	8
1	0	0	1	9
1	0	1	0	A
1	0	1	1	B
1	1	0	0	C
1	1	0	1	D
1	1	1	0	E
1	1	1	1	F

$(1100111010011010)_2$
1100 1110 1001 1010
↓ ↓ ↓ ↓
C E 9 A
 $(CE9A)_{16}$

or...

0xCE9A


Recap: Memory Management

- ❖ **Caching:** Buffering a copy of bytes (instructions and/or data) from a lower level at a higher level to exploit locality
- ❖ **Prefetching:** Preemptively retrieving bytes (typically data) from addresses not explicitly asked yet by program
- ❖ **Spill/Miss/Fault:** Data needed for program is not yet available at a higher level; need to get it from lower level
 - ❖ **Register Spill** (register to cache); **Cache Miss** (cache to main memory); “**Page**” **Fault** (main memory to disk)
- ❖ **Hit:** Data needed is already available at higher level
- ❖ **Cache Replacement Policy:** When new data needs to be loaded to higher level, which old data to evict to make room?
Many policies exist with different properties

Recap: Scheduling Policies/Algorithms


- ❖ **Schedule:** Record of **what** process runs on each CPU & **when**
- ❖ Policy controls how OS time-shares CPUs among processes
- ❖ Key terms for a process (aka **job**):
 - ❖ **Arrival Time:** Time when process gets created
 - ❖ **Job Length:** Duration of time needed for process
 - ❖ **Start Time:** Times when process first starts on processor
 - ❖ **Completion Time:** Time when process finishes/killed
 - ❖ **Response Time** = [Start Time] – [Arrival Time]
 - ❖ **Turnaround Time** = [Completion Time] – [Arrival Time]
- ❖ **Workload:** Set of processes, arrival times, and job lengths that OS Scheduler has to handle
- ❖ In general, OS may not know all Arrival Times and Job Lengths beforehand! But **preemption** is possible
- ❖ **Key Principle:** Inherent tension in scheduling between overall workload *performance* and allocation *fairness*
 - ❖ Performance metric is usually *Average Turnaround Time*

Recap: Cloud Layers




Recap: Parallelism Paradigms


Independent Workers



Shared-Nothing
Parallelism



Shared-Disk
Parallelism




Shared-Memory
Parallelism

Recap: Task Parallelism

Basic Idea: Split up *tasks* across workers; if there is a common dataset that they read, just make copies of it (aka *replication*)

Example:




This is your PA1 setup! Except, Dask Scheduler puts tasks on workers for you.

- 1) Copy whole D to all workers
- 2) Put T1 on worker 1 (W1), T2 on W2, T3 on W3; run all 3 in parallel
- 3) After T1 ends, run T4 on W1; after T2 ends, run T5 on W2; after T3 ends, W3 is *idle*
- 4) After T4 & T5 end, run T6 on W1; W2 is *idle*

Recap: Task Parallelism (continued)

- ❖ The largest amount of *concurrency* possible in the task graph, i.e., how many tasks can be run simultaneously

Example:



Q: How do we quantify the runtime performance benefits of task parallelism?

But over time, degree of parallelism keeps dropping in this example

Degree of parallelism is only 3

So, more than 3 workers is not useful for this workload!


Recap: Quantifying Parallelism Benefit

$$\text{Speedup} = \frac{\text{Completion time given only 1 worker}}{\text{Completion time given } n (>1) \text{ workers}}$$

Q: *But given n workers,
can we get a speedup of n ?*

It depends!

(On degree of parallelism, task dependency graph structure, intermediate data sizes, etc.)




Recap: Quantifying Parallelism Benefit (continued)


Scaleup refers to the ability of a system to retain the same performance ratio of tasks-per-resources when both the tasks and the resources increase at same rate.

In the above:

- "Task" can refer to a single or series of computations, queries, etc.
- "Resources" can refer to # of workers, DRAM, storage size, etc.
- "Increase" refers to using multiple instances of an initial task and initial set of resources.




Recap: Speedup vs Scaleup




Most commonly, scaling does not demonstrate ideal linear behavior.

Recap: Task Graphs and Gantt Chart

Example:



Gantt Chart visualization of schedule:



Idle times

Completion time
with 1 worker $10+5+15+5+20+10 = 65$

Parallel
completion time 35

Speedup = $65/35 = 1.9x$

Ideal/linear speedup is 3x